Correcting a bias in a climate model with an augmented emulator

Doug McNeall¹, Jonny Williams², Richard Betts^{1,3}, Ben Booth¹, and Peter Challenor³

¹Met Office Hadley Centre, FitzRoy Road, Exeter, EX1 3PB, UK

²NIWA

³University of Exeter

Correspondence: Doug McNeall (doug.mcneall@metoffice.gov.uk)

Abstract. We develop a method of bias correcting a climate model with a Gaussian process emulator, allowing valid values of

input parameters to be found even in the presence of a significant model bias.

A previous study (McNeall et al. 2016) found that a climate model had to be run using land surface input parameter values

from very different, almost non-overlapping parts of parameter space in order to satisfactorily simulate the Amazon and other

forests respectively. As the forest fraction of the other forests in the model were broadly correct at the default parameter settings

and the Amazon too low, that study suggested that the problem likely lay in the model's treatment of the Amazon region than

the other regions. The study suggested that this might be due to (1) structural errors such as missing deep-rooting in the Amazon

in the land surface component of the model, (2) a warm-dry bias in the Amazon climate of the model, or a combination of both.

In this study we bias correct the climate of the Amazon in a climate model using an "augmented" Gaussian process emulator,

where variables often regarded as outputs of the model are treated as inputs. We treat the regional temperature and precipitation

of the model as additional inputs to the emulator alongside the standard model inputs. We can then explore the relationship

between climate, input parameters and the output of the emulator, forest fraction, finding that the forest fraction is nearly as

sensitive to climate variables as any of the land surface inputs. Bias correcting the climate in the Amazon region using the

emulator corrects the forest fraction to tolerable levels in the Amazon at many candidates for land surface input parameter

values, including the default ones. It also increases the valid input space shared with that suggested by the other forests. We no

longer need to invoke a structural model error in the Amazon, beyond having too dry and hot a climate.

Using the augmented emulator allows the bias correction a pre-existing coupled ensemble of climate model runs, reducing

the risk of choosing poor parameter values because of an error in a sub-component of the model. We discuss the potential of

the augmented emulator to act as a translational layer between model sub-components simplifying the process of model tuning

when there are potential compensating errors, and helping model developers prioritise model errors to target. Our technique

has the potential to help choose good input parameters for a model, and to efficiently project the impacts of a changing climate,

even when there are significant biases in a sub-component of the model.

Copyright statement. TEXT

1

1 Introduction

Uncertainty quantification, compensating errors and model discrepancy The field of Uncertainty Quantification (UQ) has seen a rapid development of methods to quantify uncertainties when using complex computer models to simulate real, physical systems. These models often contain simplifications of processes too complex to represent explicitly in the model, termed parameterisations. Associated with these parameterisations are coefficients called input parameters, the values of which are uncertain and can be set by the model developer. The settings have a material effect on the way the parameterisations operate, and therefore on output of the model, but often to an extent that is unknown until the model is run. Input parameters are subject to uncertainty and may be difficult or even impossible to observe, having no direct analogue in the real system. The process of setting the values of the input parameters so that the simulator output best matches the real system is called tuning, and where a probability distribution is assigned for the input parameters, it is termed calibration. Uncertainty in input parameters can induce uncertainty in the output of the model, leading to uncertainty in projections of future climate states or reconstructions of past ones.

Without strong prior information it can be difficult to attribute a difference between simulator output and the real system to underlying model errors, to an incorrect set of input parameters, or to inaccuracies in the observations. Similarly, there are often a number of ways to set parameters that lead to a particular model output, with poor choices of a particular parameter compensating for poor choices of other parameters, or for modelling errors. This situation means that a good candidate for input parameters might be found in a large volume of input space, and projections of the model made with candidates from across that space might display a very wide range of outcomes. This problem is sometimes referred to as dentifiability but otherwise known as quifinality or the egeneracyf model error and parameter uncertainty. It can be relatively easy to find a good subset of input parameters given a small set of inputs and outputs and a well behaved relationship between the two. This situation might be found for a subcomponent of a climate model, where there are good observations of the system being studied, for example. Improving a coupled climate model however can require an involved and lengthy process of development. Some components of the model may have been tuned to compensate for errors in others or there may be unknown errors in the model or observations. Further, more complex models are computationally expensive and so infeasible to run in enough configurations to be able to identify these errors.

Hourdin et al. (2017) offer a summary of current practice in the somewhat understudied and sparsely documented field of climate model tuning. While there are clearly common features, there appear no standard procedures for climate model tuning however. As Hourdin et al point out, it remains an art as well as a science. Various individual centres have begun to document their tuning practices with regard to tuning targets and procedures (Schmidt et al. (2017), Zhao et al. (2018), Walters et al. (2017))

Parameter tuning occurs at different stages in model development, perhaps starting with single column version of the model. The climate model components to be coupled might be then tuned with standard boundary conditions - for example tuning a land/atmosphere component with fixed or historically observed sea surface temperatures. Finally, a system-wide tuning might be used to check that there are minimal problems once everything has been coupled together.

Golaz et al. (2013) Show the potential impact of compensating errors in tuning. They find that two different but plausible parameter configurations of the cloud formations of the coupled climate model GFDL-CM3 can result in similar present-day radiation balance. The configurations did not differ in their present day climate, but showed significantly different responses to historical forcing and therefore historical climate trajectories.

Although climate model tuning is overall a subjective process, individual parts are amenable to more algorithmic approaches. Statistical and machine learning approaches to choosing parameters to minimise modelling error, or to calculate probability distributions for parameters and model output are known as uncertainty quantification (UQ).

The problem of accounting for model discrepancy when using data to learn about input parameters is becoming more widely recognised in UQ. It was formalised in a Bayesian setting by Kennedy Oagan (2001). The authors suggested simultaneously estimating a model discrepancy - there called model inadequacy - as a function of the inputs, using a Gaussian process prior.

65

85

Arndt et al. (2012a) offer a number of examples of identifiability problems, ranging from solvable using mild assumptions through to virtually impossible. In a companion paper (Arndt et al. 2012b), they outline a way of improving identifiability using multiple model responses.

Brynjarsdottir and Oagan (2016) argued that only by accounting for model discrepancy does even a very simple simulator have a chance of making accurate predictions. Further, they found that only where there is strong prior evidence about the nature of that model discrepancy is it possible to solve the inverse problem and recover the correct inputs. Without this strong prior evidence the estimate of the correct parameters is likely to be overconfident, and wrong, leading to overconfident and wrong predictions of out-of-sample data.

Some of the dangers of overconfident and wrong estimates of input parameters and model discrepancy can be reduced using a technique called history matching (Craig et al, 1996), sometimes called pre-calibration or iterated refocussing. The aim of history matching is not to find the most likely inputs, but to reject those unlikely to produce simulations statistically close to observations of the real system. An implausibility measure (I) is calculated, taking into account the distance between the simulator output and the observation, but allowing for uncertainty in the observations, the simulator output and the simulator discrepancy. Those inputs that produce a large implausibility score are ruled out from consideration as candidate points.

An excellent introduction and case studies can be found in Andrianakis et al. (2015), or in Vernon et al. (2010). History matching is perhaps less ambitious but correspondingly more robust than calibration methods, and a full calibration can be carried out once the history matching procedure has been completed.

McNeall et al. (2013) studied an ensemble of an ice sheet model and found that using a single type of observation for ruling out input space was not very powerful - particularly if there was not a very strong relationship between an input parameter and the simulator output. A key technique therefore is to use multiple data sets for the history matching, ruling out a candidate input space according to an empirical rule. Several rules have been used - for example using the maximum implausibility of a multiple comparison, a candidate input point may be ruled out by a single observation. A more conservative approach is to use the second or third implausibility score, or to use a multivariate implausibility score, both introduced in Vernon et al. (2010). The aim of these scores is to ensure that an unidentified model discrepancy does not result in ruling out candidate points that are in fact perfectly good. History matching can be effective in reducing the volume of parameter space that is considered

plausible to produce model runs that match the real system. For example, Williamson et al. (2015) report very large reductions (around 99

95

100

120

125

While history matching has often been used used to explore and reduce the input parameter space of expensive simulators, its use as a tool to find discrepancies, bias and inadequacies in simulators is less developed. Williamson et al. (2015) argue that what was assumed a structural bias in ocean model HadCM3 could be corrected by choosing different parameters. In a different system McNeall et al. 2016 argue that a standard set of parameters for the land surface component of the climate model FAMOUS should be retained, and that a bias seen in the simulation of the Amazon rainforest is a simulator discrepancy not a poor parameter choice.

In that case, the model simulated other forests at the standard set of parameters well, and only a tiny volume of parameter space could be found that (barely) adequately simulated all the forests. When cast as a choice between keeping the default parameters, or rejecting them and accepting the new region of parameter space, they argued that the former was more likely to produce a good model, as presumably scientific judgement and expertise informed the original choice of parameters, whereas there were a number of reasons one might reject the proposed parameter space. Aims of the paper A well simulated and vigorous Amazon forest at the end of the spinup phase of a simulation experiment is a prerequisite for using the model to make robust projections of future changes in the forest. The analysis of McNeall et al. 2016 (hereafter M16) identified that the land surface input spaces where FAMOUS forest fraction was consistent with observations were very different in the Amazon than they were for other forests. The area of overlap of these spaces - one that would normally be chosen in a history matching exercise - did not simulate any of the forests well, and did not contain the default parameters. M16 suggested that assuming an error in the simulation of the Amazon forest would be a parsimonious choice. Two obvious candidates for the source of the discrepancy in the Amazon were identified: (1) a lack of deep rooting in the Amazon, meaning that trees could not access water at depth as in the real Amazon and (2) a bias in the climate of the model, impacting the vigour of the trees.

This paper revisits and extends the analysis of M16 to attempt to simultaneously (1) assess the impact of a bias corrected corrected climate on the Amazon forest and (2) to identify regions of input parameter space that should be classified as plausible, given a corrected Amazon climate. To bias correct the climate we develop a new method to augment a Gaussian process emulator, with simulator outputs acting as inputs to the emulator alongside the standard input parameters. We use simulated output of forests at different geographical locations to train the emulator, describing a single relationship between the climate of the simulator, the land surface inputs and the forest fraction. In doing so, we develop a technique that might be used to bias correct existing ensembles of coupled models, allowing a more computationally efficient method for final system-tuning of models.

In section ??, we review the literature on the possible causes of the low Amazon forest fraction in FAMOUS. In section ??, we describe how we use the temperature and precipitation to augment the Gaussian process emulator. In section ?? we use the emulator to estimate the sensitivity of forest fraction to changes in land surface and climate parameters. In section ?? we use the augmented emulator to bias correct the climates of the forest and examine the effect of that bias correction on the input space that is deemed statistically acceptable in a history matching exercise. In section ?? we search for regions of parameter space where the bias corrected simulator might perform better than at the default parameters. In section ?? we look at regions

of climate space where the default parameters would produce statistically acceptable forests. Finally, we offer some discussion
of our results in section ?? and conclusions in section ??.
2 HEADING
TEXT
2.1 HEADING
TEXT
2.1.1 HEADING
TEXT
3 Conclusions
TEXT
Code availability. TEXT
Data availability. TEXT
Code and data availability. TEXT
Sample availability. TEXT
Video supplement. TEXT

Appendix A

A1

145 Author contributions. TEXT

Competing interests. TEXT

Disclaimer. TEXT

Acknowledgements. TEXT

References

150 REFERENCE 1

REFERENCE 2